

HW3_212699581_211709597

June 5, 2022

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import numpy.linalg as lg
import numpy.random as rnd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```

Question 1:


A.

```
[2]: f = lambda x: 3 + 2 * x + x**2

lim = (-10, 10)
```

```
[3]: X = np.linspace(lim[0], lim[1], 1000)
Y = f(X)
```

```
[4]: plt.plot(X, Y, label='3 + 2x + x^2', color='red')
plt.xlabel('X')
plt.ylabel('f(X)')
plt.legend()
plt.show()
```



B.

```
[5]: grad_f = lambda x: 2 + 2 * x
```

C.

$f = 2 + 2x \setminus \begin{cases} f=0 & 2+2x=0 & x=-1 \end{cases}$

D.

```
[6]: def grad_update(grad, x, n=0.01, ep = 10**-10):
    xt_1 = x
    xt = xt_1 - n * grad(x)

    X = [xt_1, xt]

    while abs(xt - xt_1) >= ep:
        xt_1 = xt
        xt -= n * grad(xt_1)
        X.append(xt)
    return xt, np.array(X)
```

E.

```
[7]: print(grad_update(grad_f, 0, 0.01)[0])
```

-0.999999995188065

F.

```
[8]: print(len(grad_update(grad_f, 0, 0.01)[1]))
      print(len(grad_update(grad_f, 0, 0.1)[1]))
      print(len(grad_update(grad_f, -0.5, 0.01)[1]))
      print(len(grad_update(grad_f, -0.5, 0.1)[1]))

      print(len(grad_update(grad_f, 5, 0.01, 10**-6)[1]))
      print(len(grad_update(grad_f, 5, 0.1, 10**-6)[1]))
```

949

98

914

95

581

65

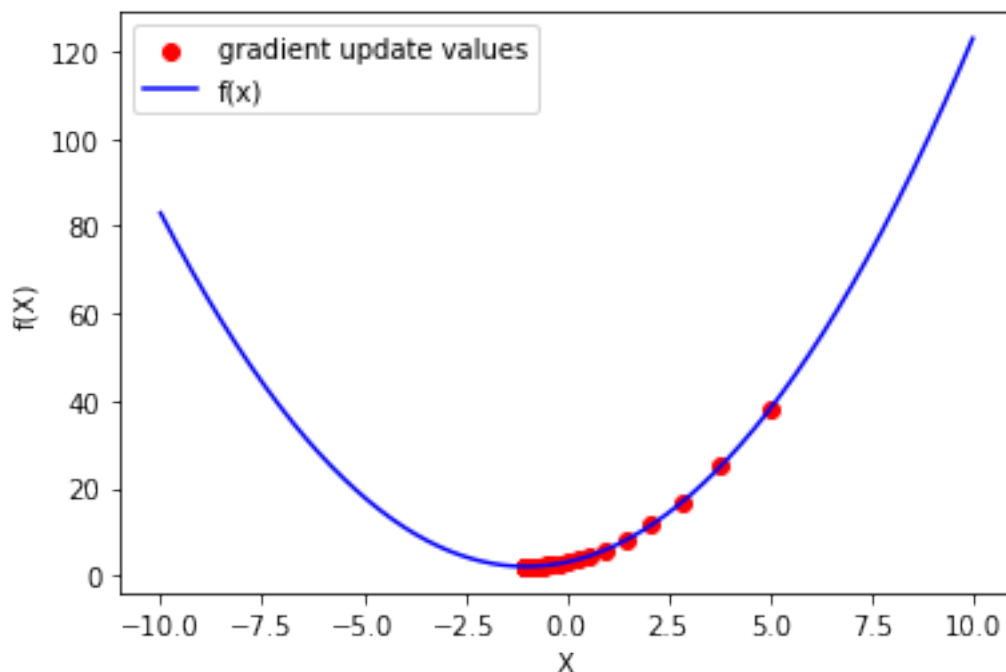
```
[9]: print(grad_update(grad_f, 5, 0.1, 10**-6)[0])
```

-0.9999962337389587

G.

```
[10]: XT = grad_update(grad_f, 5, 0.1, 10**-6)[1]
      fXT = f(XT)

      plt.scatter(XT, fXT, color='red', label='gradient update values')
      plt.plot(X, Y, color='blue', label='f(x)')
      plt.xlabel('X')
      plt.ylabel('f(X)')
      plt.legend()
      plt.show()
```



Question 2:

A:

As we've seen in lectures $\lambda\|w\|^2$ is 2λ -strongly convex

And we've seen that $\max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}$ is convex, and from a theorem from algebra if f and g is convex then $f+g$ is convex as well and so $\frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\} + \lambda\|w\|^2$ is convex

B:

we need to proof that $|l(w_1, x_i, y_i) - l(w_2, x_i, y_i)| < R\|w_1 - w_2\|$

- $|l(w_1, x_i, y_i) - l(w_2, x_i, y_i)| =$
 $|\max\{0, 1 - y_i \langle w_1, x_i \rangle\} - \max\{0, 1 - y_i \langle w_2, x_i \rangle\}| =$
 - i. $1 - y_i \langle w_1, x_i \rangle - 0 \leq 1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_2, x_i \rangle) = \max\{0, 1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_1, x_i \rangle)\}$
 - ii. $1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_1, x_i \rangle) = \max\{0, 1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_1, x_i \rangle)\}$
 - iii. $0 - (1 - y_i \langle w_2, x_i \rangle) \leq 0 = \max\{0, 1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_1, x_i \rangle)\}$
 - iv. $0 - 0 \leq \max\{0, 1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_1, x_i \rangle)\}$

** $i. \max_k \|x_k\| \|w_2 - w_1\| \geq \|x_i\| \|w_2 - w_1\|$ cauchy schwarz inequality

$\langle x_i, w_2 - w_1 \rangle = \langle x_i, w_2 \rangle - \langle x_i, w_1 \rangle =$

$-1 + \langle x_i, w_2 \rangle + 1 - \langle x_i, w_1 \rangle =$

a. $1 - y_i \langle x_i, w_1 \rangle - (1 - y_i \langle x_i, w_2 \rangle)$, for $y_i = 1$

b. for $y_i = -1$ start with $\|w_1 - w_2\|$ and we'll get the same inequality

ii. $\|x_i\| \|w_2 - w_1\| \geq 0$

$\Rightarrow \max_k \|x_k\| \|w_2 - w_1\| \geq \max\{0, 1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_2, x_i \rangle)\}$

combine * and ** and we get $|l(w_1, x_i, y_i) - l(w_2, x_i, y_i)| < R \|w_1 - w_2\|$

C:

Denote $\mathcal{L}(w, b) = \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\} + \lambda \|w\|^2$

$$\frac{\partial \mathcal{L}}{\partial w} = \begin{cases} 2\lambda w, & \text{if } 1 - y_i(\langle w, x_i \rangle + b) \leq 0 \\ -y_i x_i + 2\lambda w, & \text{o.w.} \end{cases}$$

$$\frac{\partial \mathcal{L}}{\partial b} = \begin{cases} 0, & \text{if } 1 - y_i(\langle w, x_i \rangle + b) \leq 0 \\ -y_i, & \text{o.w.} \end{cases}$$

D:

```
[11]: def sub_grad(w, b, x, y, lam, d, m):
    return 2*lam*w if (1 - y * (w@x + b)) <= 0 else -y*x+2*lam*w, 0 if (1 - y *
    ↪(w@x + b)) <= 0 else -y

def pract_sgd(X, y, lam, epochs, l_rate):
    m = len(X)
    d = len(X[0])

    w = rnd.uniform(0, 1, d)
    b = rnd.uniform(0, 1)

    perm = np.arange(m)

    for i in range(epochs):
#         print(f'epoch {i}. {float(i)/epochs:.2f}%')
        rnd.shuffle(perm)
        for i in perm:
            subGrad = sub_grad(w, b, X[i], y[i], lam, d, m)
            w -= l_rate * subGrad[0]
            b -= l_rate * subGrad[1]

    return w, b

def theory_sgd(X, y, lam, epochs, l_rate):
    m = len(X)
    d = len(X[0])
```

```

w = rnd.uniform(0, 1, d)
b = rnd.uniform(0, 1, 1)

W = []
B = []

W.append(w)
B.append(b)

for i in range(m*epochs):
#     print(f'epoch {i}. {float(i)/(m*epochs):.2f}%')
    index = rnd.randint(0, m)
    subGrad = sub_grad(w, b, X[index], y[index], lam, d, m)
    w -= l_rate * subGrad[0]
    b -= l_rate * subGrad[1]
    W.append(w)
    B.append(b)

return sum(W) / (m * epochs), sum(b) / (m * epochs)

def svm_with_sgd(X, y, lam=0, epochs=1000, l_rate=0.01, sgd_type="practical"):
    rnd.seed(2)
    if sgd_type == "practical":
        return pract_sgd(X, y, lam, epochs, l_rate)
    return theory_sgd(X, y, lam, epochs, l_rate)

```

E:

```

[12]: sign = lambda x: 1 if x >= 0 else -1

def calculate_error(w, b, X, y):
    c = 0
    for x_i, y_i in zip(X, y):
        if sign(w*x_i + b) != y_i:
            c += 1
    return c / len(y)

```

F:

```
[13]: X, y = load_iris(return_X_y=True)
X = X[y != 0]
y = y[y != 0]
y[y == 2] = -1
X = X[:, 2:4]

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3,
↪random_state=0)

Lam = [0, 0.05, 0.1, 0.2, 0.5]

train_err = list()
val_err = list()
margin = list()

for l in Lam:
    w, b = svm_with_sgd(X_train, y_train, l)
    train_err.append(calculate_error(w, b, X_train, y_train))
    val_err.append(calculate_error(w, b, X_val, y_val))
    margin.append(1/lg.norm(w))
```

```
[14]: fig, ax = plt.subplots(2, 1)

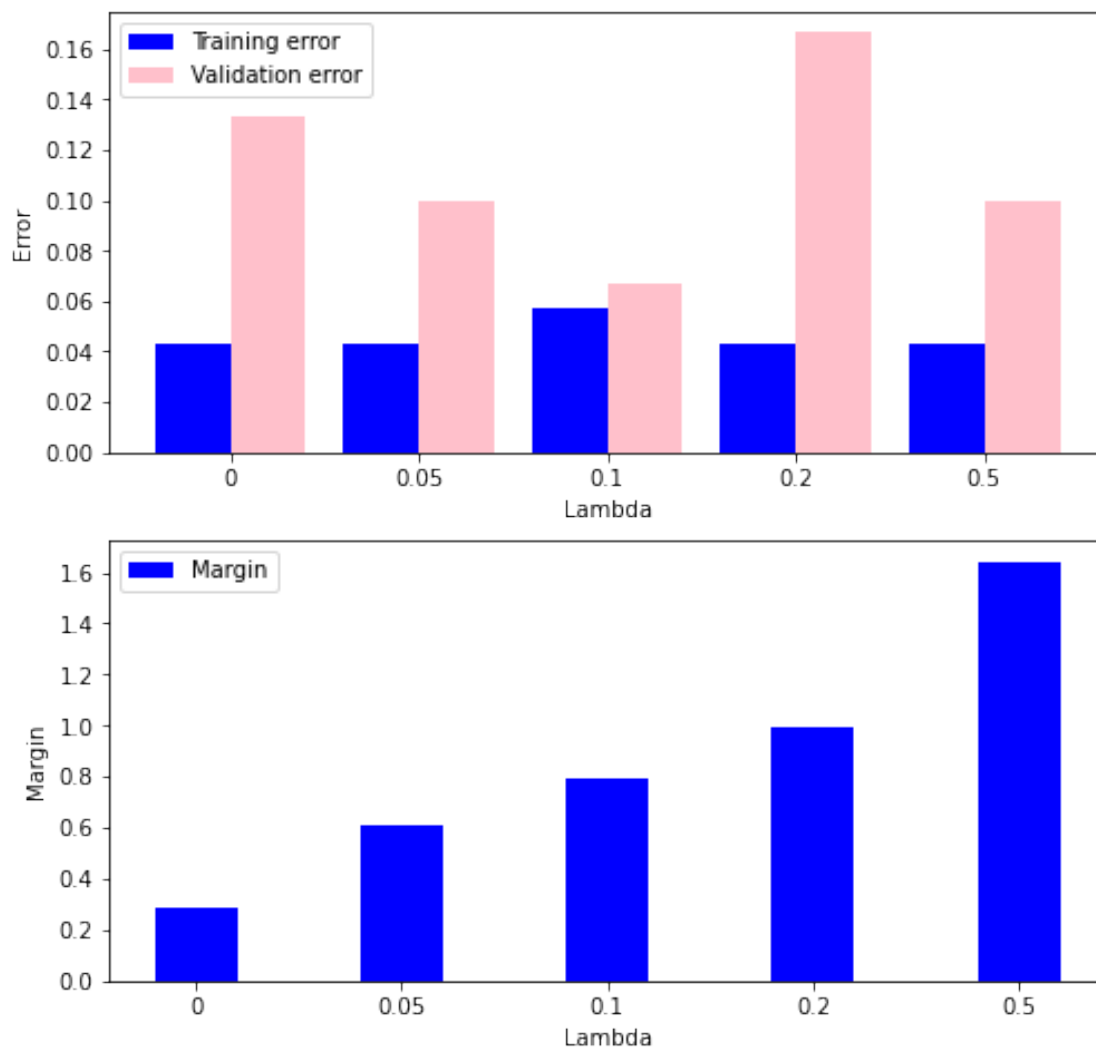
fig.set_size_inches(8, 8)

x = np.arange(len(Lam))

ax[0].bar(x-0.4/2, train_err, 0.4, color="blue", label="Training error")
ax[0].bar(x+0.4/2, val_err, 0.4, color="pink", label="Validation error")
ax[0].set_xticks(x)
ax[0].set_xticklabels(Lam)
ax[0].set_xlabel("Lambda")
ax[0].set_ylabel("Error")
ax[0].legend()

ax[1].bar(x, margin, 0.4, color="blue", label="Margin")
ax[1].set_xticks(x)
ax[1].set_xticklabels(Lam)
ax[1].set_xlabel("Lambda")
ax[1].set_ylabel("Margin")
ax[1].legend()

plt.show()
```



We see that with $\lambda = 0.1$ we get the best validation error

G:

```
[20]: Epochs = np.arange(10, 1001, 10)
practical_train_error = list()
theoretical_train_error = list()

practical_validation_error = list()
theoretical_validation_error = list()

for epoch in Epochs:
    if epoch % 200 == 0:
        print(f'epoch {epoch}')
```



```

w1, b1 = svm_with_sgd(X_train, y_train, lam=0.1, epochs=epoch,
↪sgd_type="practical")
w2, b2 = svm_with_sgd(X_train, y_train, lam=0.1, epochs=epoch,
↪sgd_type="theory")
practical_train_error.append(calculate_error(w1, b1, X_train, y_train))
theoretical_train_error.append(calculate_error(w2, b2, X_train, y_train))

practical_validation_error.append(calculate_error(w1, b1, X_val, y_val))
theoretical_validation_error.append(calculate_error(w2, b2, X_val, y_val))

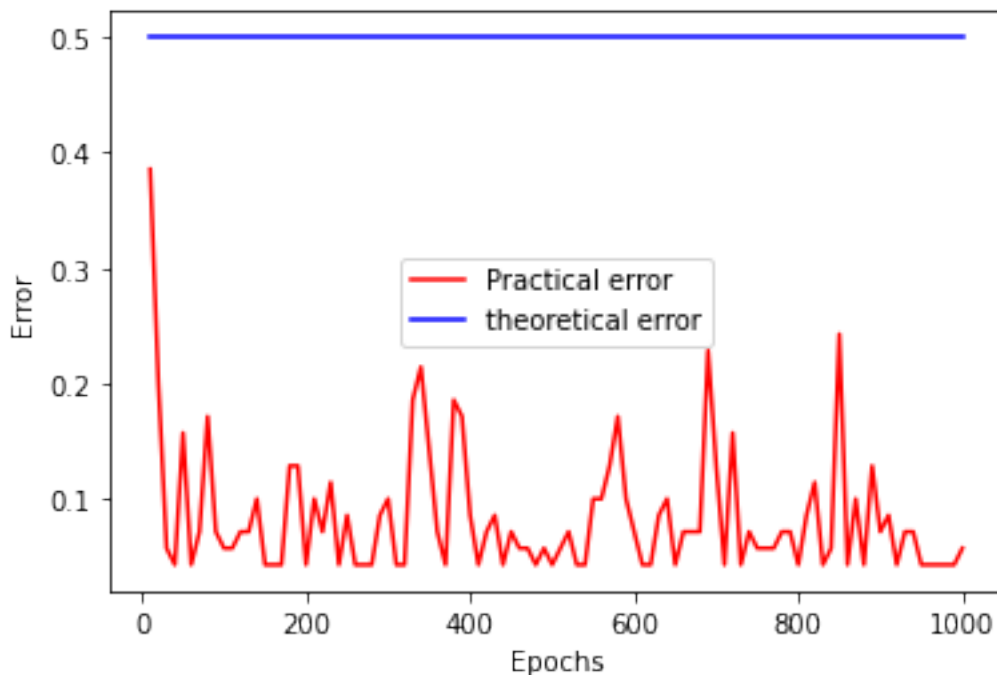
```

```

[18]: plt.plot(Epochs, practical_train_error, color="red", label="Practical error")
plt.plot(Epochs, theoretical_train_error, color="blue", label="theoretical_
↪error")
plt.xlabel("Epochs")
plt.ylabel("Error")
plt.legend()
plt.plot()

```

[18]: []



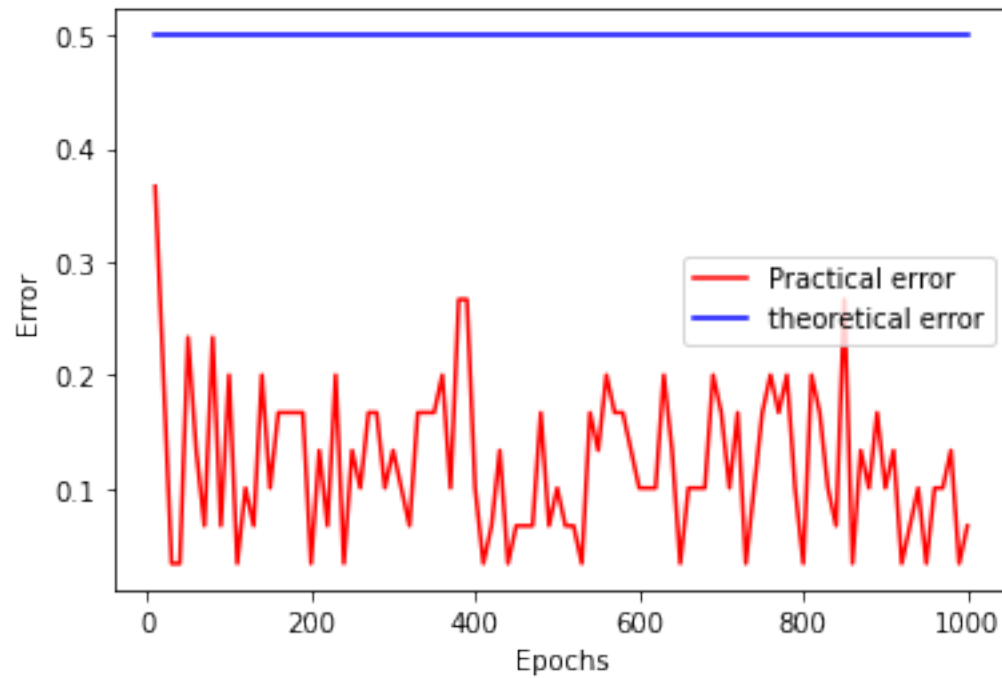
```

[19]: plt.plot(Epochs, practical_validation_error, color="red", label="Practical_
↪error")

```

```
plt.plot(Epochs, theoretical_validation_error, color="blue", label="theoretical_
↪error")
plt.xlabel("Epochs")
plt.ylabel("Error")
plt.legend()
plt.plot()
```

[19]: []



: 3 = 8k e

```
#Q3_A:
import numpy as np
#folds=k
def cross_validation_error(X, y, model, folds):
    X_k=np.array_split(X, folds)
    Y_k=np.array_split(y, folds)
    train_error=[]
    val_error=[]
    for fold_1 in range(folds):
        X_val_K=X_k[fold_1]
        Y_val_K=Y_k[fold_1]
        X_train_K=np.array([element for fold_2 in range(folds) if fold_2!=fold_1 for element in X_k[fold_1]])
        Y_train_K=np.array([element for fold_2 in range(folds) if fold_2!=fold_1 for element in Y_k[fold_1]])
        model.fit(X_train_K,Y_train_K)
        train_error.append(model.score(X_train_K,Y_train_K))
        val_error.append(model.score(X_val_K,Y_val_K))
    return(1-(np.array(train_error).mean()),1-(np.array(val_error).mean()))
```

-1e

```
[2] #Q3_B:
from sklearn.svm import SVC
def svm_results(X_train, y_train, X_test, y_test):
    folds=5
    lamda=[0.0001, 0.01, 1, 100, 10000 ]
    output={}
    for lamda in lamda:
        model=SVC(kernel='linear',C=1/lamda)
        train_error,val_error=cross_validation_error(X_train, y_train, model, folds)
        model.fit(X_train,y_train)
        test_error= (model.predict(X_test)!=y_test).mean()
        output[f'svm_lambda_{lamda}']=(train_error,val_error,test_error)
    return output
```

-2

```
#Q3_C:
from sklearn.datasets import load_iris
iris_data = load_iris()
X, y = iris_data['data'], iris_data['target']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=7)
output=svm_results(X_train, y_train, X_test, y_test)
```

-2

```

import numpy as np
import matplotlib.pyplot as plt

# set width of bar
barWidth = 0.25
fig = plt.subplots(figsize =(12,8))

# set height of bar
train=[]
val=[]
test=[]
for i in range(len(list(output.values()))):
    train.append(list(output.values())[i][0])
    test.append(list(output.values())[i][2])
    val.append(list(output.values())[i][1])


# Set position of bar on X axis
br1 = np.arange(len(train))
br2 = [x + barWidth for x in br1]
br3 = [x + barWidth for x in br2]

# Make the plot
plt.bar(br1, train, width = barWidth,
        edgecolor = 'black', label = 'train_Error')
plt.bar(br2, val, width = barWidth,
        edgecolor = 'black', label = 'val_Error')
plt.bar(br3, test, width = barWidth,
        edgecolor = 'black', label = 'test_Error')

# Adding Xticks
plt.xlabel('lamda', fontweight = 'bold', fontsize = 15)
plt.ylabel('Errors', fontweight = 'bold', fontsize = 15)
plt.xticks([r + barWidth for r in range(len(train))],
           [0.0001, 0.01, 1, 100, 10000 ])

plt.legend()
plt.show()

```



הוסבר :

המדד האם ביותר דפי זיגור λ עם מדד הדבר
הוא עבור $\lambda=1$ כי הליגיה ($train$ error) היא מניח
בדיוק כלל

אני מחפשים הליגיה עם ה $train$ error & val error
אזי מקיים עכ"ר $\lambda=1$ אכן עבור ערך ה λ
כלי מקבל המדד האם ביותר דפי זיגור λ .

נניח אם מדד עכ"ר ה קטנים ($\lambda < 1$) יש ליה
זכור ה $overfitting$ כי בערכים האלה $train$ error
זכור ה מדד ביחס ל val error.
במקרה עכ"ר ה גדולים הם זכור ה $train/test/val$
כלומר יש ליה $underfitting$.

אכן בליגיה האחרונה המדד האם ביותר היא
עם $\lambda=1$.

שאלה 4:

$$\forall i \in \{1, \dots, r\} \quad g_i: \mathbb{R}^d \rightarrow \mathbb{R}$$

כך שהפונקציה קמורה ולכיתה בכל \mathbb{R}^d .

$$g(w) = \max_{i \in [r]} g_i(w)$$

נאכיח עכשיו גרסאות מסוימות של

$$j \in \arg \max_i g_i(w)$$

ע $\nabla g_j(w)$ סגור-לכיתה על הפונקציה g הנקודה w .

כגוף: $\frac{\cdot}{\cdot}$
לכל $v \in \mathbb{R}^d$ מקיים:

$$g(w) + \langle v - w, \nabla g_j(w) \rangle \underset{\substack{\text{מהצורה } g}}{=} g_j(w) + \langle v - w, \nabla g_j(w) \rangle$$

נכון ע g_j פונקציה קמורה ולכיתה בכל \mathbb{R}^d ובכך
בהנחה w חסר דפי סף שוויון כזו קיימת:

$$\forall u \in \mathbb{R}^d: \quad g_j(u) \geq g_j(w) + \langle v - w, \nabla g_j(w) \rangle$$

ואם דפי הצורה הפונקציה: $g_j(u) = g(u)$
 $g_j(u) \leq \max_{i \in [r]} g_i(u) = g(u)$ \square